

02 Back Up and Restore with MongoDB Tools

The `mongodump` and `mongorestore` utilities work with BSON data dumps, and are useful for creating backups of small deployments.

When backing up your data with MongoDB's tools, consider the following guidelines:

- Label files so that you can identify the contents of the backup as well as the point in time that the backup reflects.
- Use an alternative backup strategy such as Filesystem Snapshots or MongoDB Cloud Manager if the performance impact of `mongodump` and `mongorestore` is unacceptable for your use case.
- Use `--oplog` to capture incoming write operations during the `mongodump` operation to ensure that the backups reflect a consistent data state.
- Ensure that your backups are usable by restoring them to a test MongoDB deployment.

Basic mongodump Operations

The `mongodump` utility backs up data by connecting to a running `mongod`. The utility can create a backup for an entire server, database or collection, or can use a query to backup just part of a collection.

When you run `mongodump` without any arguments, the command connects to the MongoDB instance on the local system (e.g. `localhost`) on port 27017 and creates a database backup named `dump/` in the current directory.

Sintaxis

```
mongodump --host=<dominio-o-ip> --port=27017
```

To specify a different output directory, you can use the `--out` or `-o` option:

```
mongodump ... --out=/data/backup/
```

Práctica

Creamos dos directorios

```
md data/server
```

```
md data/backup
```

y levantamos un servidor:

```
mongod --port 27100 --dbpath data/server
```

Escribimos:

```
> use getafe1
switched to db getafe1
> for (i=0; i < 1000; i++) {db.foo1.insert({a: i})}
WriteResult({ "nInserted" : 1 })
> use getafe2
switched to db getafe2
> for (i=0; i < 1000; i++) {db.foo2.insert({a: i})}
WriteResult({ "nInserted" : 1 })
```

Desde otra terminal lanzamos mongodump:

```
MacBook-Pro-de-Pedro:~ pedro$ mongodump --port 27100
--out=data/backup/mongodump-11-02-2020
2020-02-11T17:23:49.095+0100    writing admin.system.version to
2020-02-11T17:23:49.096+0100    done dumping admin.system.version (1
document)
2020-02-11T17:23:49.097+0100    writing getafe1.foo1 to
2020-02-11T17:23:49.098+0100    writing getafe2.foo2 to
2020-02-11T17:23:49.099+0100    done dumping getafe1.foo1 (1000 documents)
2020-02-11T17:23:49.100+0100    done dumping getafe2.foo2 (1000 documents)
```

To limit the amount of data included in the database dump, you can specify `--db` and `--collection` as options to `mongodump`. For example:

```
mongodump ... --collection=<colección> --db=<basededatos>
```

`mongodump` overwrites output files if they exist in the backup data folder. Before running the `mongodump` command multiple times, either ensure that you no longer need the files in the output folder (the default is the `dump/` folder) or rename the folders or files.

Point in Time Operation Using Oplogs

Use the `--oplog` option with `mongodump` to collect the oplog entries to build a point-in-time snapshot of a database within a replica set. With `--oplog`, `mongodump` copies all the data from the source database as well as all of the oplog entries from the beginning to the end of the backup procedure. This

operation, in conjunction with mongorestore --oplogReplay, allows you to restore a backup that reflects the specific moment in time that corresponds to when mongodump completed creating the dump file.

Restore a Database with mongorestore

The mongorestore utility restores a binary backup created by mongodump. By default, mongorestore looks for a database backup in the dump/ directory.

The mongorestore utility restores data by connecting to a running mongod directly.

mongorestore can restore either an entire database backup or a subset of the backup.

To use mongorestore to connect to an active mongod, use a command with the following prototype form:

Sintaxis

```
mongorestore --port=<numero-puerto> <ruta>
```

Práctica

Borramos una de las colecciones:

```
> db.foo2.drop()
true
```

Y restauramos:

```
MacBook-Pro-de-Pedro:~ pedro$ mongorestore --port 27100
data/backup/mongodump-11-02-2020
```

```
...
```

```
2020-02-11T18:31:52.058+0100    no indexes to restore
2020-02-11T18:31:52.058+0100    finished restoring getafe1.foo1 (0 documents, 1000
failures)
2020-02-11T18:31:52.058+0100    1001 document(s) restored successfully. 1000
document(s) failed to restore.
```

Ahora comprobamos como está restituida la colección borrada.

```
> db.foo2.find()
{ "_id" : ObjectId("5e42d3864fac851c55bc1337"), "a" : 0 }
{ "_id" : ObjectId("5e42d3864fac851c55bc1338"), "a" : 1 }
```

```
{ "_id" : ObjectId("5e42d3864fac851c55bc1339"), "a" : 2 }  
{ "_id" : ObjectId("5e42d3864fac851c55bc133a"), "a" : 3 }  
{ "_id" : ObjectId("5e42d3864fac851c55bc133b"), "a" : 4 }  
...
```